

CGC: a scalable Python package for co- and tri-clustering of geo-data cubes

18 November 2021

Abstract

Data cubes, i.e., datasets whose variables are expressed as a function of multiple dimensions (e.g., space, time, band), are ubiquitous in the field of Earth observation and geoscience. Clustering techniques encompassing the full dimensionality of the data are necessary to identify patterns that extend across multiple axes in such datasets. Clustering Geodata Cubes (CGC) is a Python package designed to enable this class of analyses, especially targeting multi-dimensional geospatial datasets. It provides the functionality to perform co- and tri-cluster analysis, offering scalable implementations that suit both small and large datasets.

Introduction

Faced with the increasing ubiquity of large datasets, data mining techniques have become essential to extracting patterns and generating insights. In this regard, clustering techniques, which aim to identify groups or subgroups with similar properties within a larger data set, are becoming ever more popular.

In the case of multi-variate data, however, traditional one-dimensional clustering techniques, which marginalize over the other variables, may obfuscate clusters present in the data (Hartigan 1972; Cheng and Church 2000). Therefore, techniques capable of grouping data along multiple axes simultaneously are needed. These techniques are referred to as co- or bi-clustering in the case of two axes or tri-clustering in the case of three axes and have seen significant development and adoption in fields ranging from bioinformatics (Cheng and Church 2000) to finance (Shi et al. 2018) and natural language processing (Dhillon 2001).

With many large datasets referring to geographical phenomena and either describing time-varying properties (so-called GTS - georeferenced time series, including, e.g. satellite image time series) or providing multi-dimensional information about a given location or phenomenon (e.g. in weather and climate datasets),

geo-information science is a further natural field of employment for co- and tri-clustering. Indeed, following the development of a general information-theoretical approach (Dhillon, Mallela, and Modha 2003) to partitional co-clustering (Banerjee et al. 2007), Wu et al. presented an application of co-clustering to geo-information science (Wu, Zurita-Milla, and Kraak 2015), as well as an extension of the method to tri-clustering (Wu et al. 2018). As they argue, with the explosion of Earth observation and remote sensing, and therefore GTS in three dimensions, such methods will become increasingly essential in tackling the large data volumes becoming available.

In light of the eminent employability of partitional co- and tri-clustering to geo-information science and Earth observation, but also the transferability to other domains, this paper presents and publishes the implementation in the Clustering Geo-data Cubes (CGC) package of the co-clustering and tri-clustering methods originally developed in Wu, Zurita-Milla, and Kraak (2015) and Wu et al. (2018), respectively.

As outlined below, the package aims to meet the needs of the geo-information science community, including, in particular, the ability to handle ‘big data’ and to readily interface with the existing geographic information system (GIS) software ecosystem. Nevertheless, the methods remain widely applicable and can easily be applied in other domains as well.

Statement of need

The CGC package focuses on the needs of geographers and geoscientist. In particular, it aims to meet the community need for a tool that can accurately cluster multi-dimensional data by providing the following features and functionalities:

- **Partitional co- and tri-clustering algorithms, as suitable to work with spatiotemporal (multi-dimensional) data.** CGC entails algorithms that are based on information theory and that are designed to simultaneously group elements into disjoint clusters along all the input array axes. These methods are advantageous over single-sided clustering in that they provide a strategy to identify patterns that unfold across multiple dimensions (e.g. space and time). In addition, CGC provides the functionality to perform a cluster refinement over the identified co- and tri-clusters. This post-processing step has the goal to reduce the number of clusters by grouping them based on similarity, facilitating the identification and visualization of patterns.
- **A scalable approach to efficiently perform cluster analyses on both small and big data sets.** CGC offers solutions to make efficient use of the available computational resources and to tackle a wide range of data set sizes. It provides co- and tri-clustering implementations designed to run on either a single machine or on a compute cluster. For the former, which tackles input data arrays that largely fit into memory, one can exploit

parallelization of independent analysis runs, taking advantage of multi-core CPUs. For the latter, which tackles instead large datasets, the analysis is carried out using distributed data and computation.

- **A framework easy to integrate into geospatial analysis workflows.**

CGC is written in Python, which is one of the top programming languages for geospatial scripting and applications. In the implementation targeting distributed computing, CGC makes use of the Dask library (Dask Development Team 2016), which is widely employed in the field of big geo-data. Numpy and Dask arrays, which are common data structures in Python, are employed as input and output data types in CGC. These structures are also employed in the higher-level objects in the Xarray package (Hoyer and Hamman 2017), so that this versatile and popular tool can be used for data loading and manipulation before and after analyses with CGC. Documentation and tutorials illustrate domain-science examples, applications, and use cases to facilitate community adoption. To make CGC findable and easy to install and to provide a platform for issue tracking, the development takes place in a publicly accessible repository, the package is distributed via the Python Package Index (PyPI), and code-release snapshots are archived on Zenodo. Software quality and reliability are guaranteed via unit tests implemented in the continuous integration cycle.

Algorithms

Co-clustering

CGC implements the Bregman block average co-clustering (BBAC) algorithm from Banerjee et al. (2007) (see Algorithm 1 in the article), and it was inspired by the Matlab code by Srujana Merugu and Arindam Banerjee (Merugu and Banerjee 2004). The algorithm iteratively optimizes the clustering of rows and columns starting from a random initial assignment. The information loss from the original matrix to the clustered one, which is constructed as the matrix of the co-cluster means, is minimized using a loss function that is based on the I-divergence. The algorithm is considered converged when the difference in loss function between two consecutive optimization steps is smaller than a provided threshold. To limit the influence of the initial conditions on the final clustering, which might represent a local minimum in the cluster space, multiple differently-initialized runs are carried out. Ultimately, the cluster assignment is selected from the run that leads to the lowest loss-function value.

The iterative cluster optimization involves steps where the cluster-based means are calculated and the row- and column-cluster assignments updated to minimize the loss function. Note that in the CGC implementation of the algorithm, the update in the row- and column-cluster assignments is computed only from the previous iteration’s row and column clusters (and the corresponding cluster-based means). Contrarily to the original MATLAB implementation (Merugu and Banerjee 2004), this makes the algorithm independent of the order in which

the dimensions are considered, while still leading to a (locally) optimal clustering solution.

Tri-clustering

For tri-clustering, CGC implements the so-called Bregman cube average tri-clustering (BCAT) algorithm, which is the natural generalization of the BBAC algorithm to three dimensions (Wu et al. 2018). A loss function based on the I-divergence is employed and multiple runs are carried out to avoid local minima as much as possible. At every iteration, the cluster-based means are calculated and the clusters in the three dimensions updated to minimize the information loss from the original data matrix to its clustered counterpart. Similar to the CGC co-clustering implementation, only the clusters from the previous iteration (and the corresponding cluster means) are employed to update the tri-cluster assignments. Thus, the tri-clustering outcome is not influenced by the order in which the three axes of the input data array are provided.

Cluster refinement

The CGC package implements an optional cluster refinement step based on the k-means method (Wu, Zurita-Milla, and Kraak 2016). For this, we exploit the k-means implementation available in the scikit-learn package (Pedregosa et al. 2011). The co- and tri-clusters are grouped into k pre-defined clusters. This grouping is based on statistical properties of the co- or tri-clusters and helps to better capture the patterns hidden in the data. The statistical properties employed by CGC are listed in the package documentation. The optimal k value is selected within a given range, using the Silhouette metric (Rousseeuw 1987).

Software package overview

The CGC software is structured in the following main modules, whose details are described in the online package documentation:

- **coclustering**, which offers access to all the available implementations of the same co-clustering algorithm:
 - An implementation based on Numpy, as suitable to run the algorithm on a single machine. In this implementation, vectorized operations between work arrays are exploited to efficiently calculate cluster-based means and to optimize the cluster assignments. Multiple differently-initialized runs can be run in parallel using the threading library, allowing to make optimal usage of machines with multi-core CPUs.
 - A second Numpy-based implementation that replaces some of the vectorized operations with looping to reduce the need for large work arrays, leading to a much-reduced memory footprint (negligible with respect to the size of the data matrix). While the reduced

memory requirement comes at the cost of performance (“for” loops in Python are slow), the loss can be mitigated through a feature that activates Numba’s just-in-time compilation for some of the critical steps.

- An implementation based on the Dask library, targeting data sets whose size would prevent analyses on a single machine. In particular, Dask arrays are employed to process the data in chunks, which are distributed across the nodes of a compute cluster. The runs to sample initial conditions are executed sequentially to keep memory requirements to the minimum.
- **triclustering**, which offers access to all tri-clustering implementations:
 - A Numpy-based implementation analogous to the co-clustering one described above (note that the low-memory version is currently not available).
 - A Dask-based implementation, also analogous to the corresponding co-clustering version described above.
- **kmeans**, which implements the k-means cluster refinement step for both co- and tri-clustering.
- **utils**, which includes a collection of utility functions e.g. memory consumption estimation and cluster averaging.

Tutorial

The software package is complemented by an online tutorial that illustrates how to perform cluster analysis of geospatial datasets using CGC. The tutorial exploits a format that is widely employed in the geoscience community, i.e. Jupyter notebooks (Kluyver et al. 2016). Notebooks are made directly available via a dedicated GitHub repository, but they are also published as static web pages for reference and linked to the CGC documentation via a ‘Tutorials’ tab. The tutorials are designed to run on systems with limited CPU/memory capacity, which, together with environment requirement specifications in a standardized format (conda YAML file) and binder badges, give users the possibility to easily run the notebooks live on mybinder.org (Jupyter Project et al. 2018).

Tutorials cover the following topics:

- Co- and tri-cluster analysis.
- K-means-based cluster refinement.
- Choice of the suitable implementation for a given problem size/infrastructure available.
- Loading of geospatial data, common data-manipulation tasks in a workflow involving CGC, visualization of the output.

Note that while the tutorial is aimed at geospatial use cases, it illustrates some

real-case applications that are likely to make it easier for users to carry out cluster analysis using CGC in other fields as well.

Acknowledgements

The authors would like to thank Dr. Yifat Dzigan for the helpful discussions and support and Dr. Romulo Goncalves for the preliminary work that lead to the development of the software package presented here. We also would like to thank SURF for providing computational resources to test the first versions of the CGC package via the EINF-XXXX grant.

Author contributions

All co-authors contributed to the conceptualization of the work, which was led by R.Z.M. and E.I.V.. F.N., M.W.G., and O.K. prepared the first draft of the tutorials, and wrote the initial draft of this manuscript. F.N. suggested changes to the co- and tri-clustering algorithms. R.Z.M., E.I.V., and S.K. led the design of experiments to test and improve the CGC package. R.Z.M. and E.I.V. helped to improve the tutorials. S.K. provided the required computational resources to run the experiments and tutorials and made suggestions for code optimizations. All co-authors reviewed and edited the final document.

References

- Banerjee, Arindam, Inderjit Dhillon, Joydeep Ghosh, Srujana Merugu, and Dharmendra S. Modha. 2007. “A Generalized Maximum Entropy Approach to Bregman Co-Clustering and Matrix Approximation.” *The Journal of Machine Learning Research* 8 (December): 1919–86.
- Cheng, Yizong, and George M Church. 2000. “Biclustering of Expression Data.” In *Ismb*, 8:93–103. 2000.
- Dask Development Team. 2016. *Dask: Library for Dynamic Task Scheduling*. <https://dask.org>.
- Dhillon, Inderjit S. 2001. “Co-Clustering Documents and Words Using Bipartite Spectral Graph Partitioning.” In *Proceedings of the Seventh ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 269–74. KDD ’01. New York, NY, USA: Association for Computing Machinery. <https://doi.org/10.1145/502512.502550>.
- Dhillon, Inderjit S., Subramanyam Mallela, and Dharmendra S. Modha. 2003. “Information-Theoretic Co-Clustering.” In *Proceedings of the Ninth ACM SIGKDD International Conference on Knowledge Discovery and Data Mining - KDD ’03*, 89. Washington, D.C.: ACM Press. <https://doi.org/10.1145/956750.956764>.
- Hartigan, J. A. 1972. “Direct Clustering of a Data Matrix.” *Journal of the American Statistical Association* 67 (337): 123–29. <https://doi.org/10.1080/>

01621459.1972.10481214.

- Hoyer, Stephan, and Joseph J. Hamman. 2017. "Xarray: N-D Labeled Arrays and Datasets in Python." *Journal of Open Research Software* 5 (April): 10. <https://doi.org/10.5334/jors.148>.
- Jupyter Project, Matthias Bussonnier, Jessica Forde, Jeremy Freeman, Brian Granger, Tim Head, Chris Holdgraf, et al. 2018. "Binder 2.0 - Reproducible, Interactive, Sharable Environments for Science at Scale." In, 113–20. Austin, Texas. <https://doi.org/10.25080/Majora-4af1f417-011>.
- Kluyver, Thomas, Benjamin Ragan-Kelley, Fernando Pérez, Brian Granger, Matthias Bussonnier, Jonathan Frederic, Kyle Kelley, et al. 2016. "Jupyter Notebooks - a Publishing Format for Reproducible Computational Workflows." In *Positioning and Power in Academic Publishing: Players, Agents and Agendas*, edited by Fernando Loizides and Birgit Schmidt, 87–90. IOS Press. <https://eprints.soton.ac.uk/403913/>.
- Merugu, Srujana, and Arindam Banerjee. 2004. "Bregman Co-Clustering Code in Matlab." <http://www.ideal.ece.utexas.edu/software.html>.
- Pedregosa, F., G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, et al. 2011. "Scikit-Learn: Machine Learning in Python." *Journal of Machine Learning Research* 12: 2825–30.
- Rousseeuw, Peter J. 1987. "Silhouettes: A Graphical Aid to the Interpretation and Validation of Cluster Analysis." *Journal of Computational and Applied Mathematics* 20 (November): 53–65. [https://doi.org/10.1016/0377-0427\(87\)90125-7](https://doi.org/10.1016/0377-0427(87)90125-7).
- Shi, Guangwei, Liying Ren, Zhongchen Miao, Jian Gao, Yanzhe Che, and Jidong Lu. 2018. "Discovering the Trading Pattern of Financial Market Participants: Comparison of Two Co-Clustering Methods." *IEEE Access* 6: 14431–38.
- Wu, Xiaojing, Raul Zurita-Milla, Emma Izquierdo-Verdiguier, and Menno-Jan Kraak. 2018. "Triclustering Georeferenced Time Series for Analyzing Patterns of Intra-Annual Variability in Temperature." *Annals of the American Association of Geographers* 108 (1): 71–87. <https://doi.org/10.1080/24694452.2017.1325725>.
- Wu, Xiaojing, Raul Zurita-Milla, and Menno-Jan Kraak. 2015. "Co-Clustering Geo-Referenced Time Series: Exploring Spatio-Temporal Patterns in Dutch Temperature Data." *International Journal of Geographical Information Science* 29 (4): 624–42. <https://doi.org/10.1080/13658816.2014.994520>.
- . 2016. "A Novel Analysis of Spring Phenological Patterns over Europe Based on Co-Clustering: Co-Clustering European Spring Phenology." *Journal of Geophysical Research: Biogeosciences* 121 (6): 1434–48. <https://doi.org/10.1002/2015JG003308>.